

Group 7: GIL-free Python

Code: github.com/sueszli/nogil

In October 2024, for the first time Python overtook JavaScript as the most popular language on Github’s Octoverse¹. This is a testament to Python’s versatility and ease of use, making it an ideal choice for scripting and prototyping. However, the language’s simplicity comes at a cost: It is notoriously slow for compute-bound tasks due to the Global Interpreter Lock (GIL)². This lock (in the most popular implementation, CPython) prevents multiple native threads from executing Python bytecodes simultaneously, effectively limiting the language’s performance on multi-core systems³. This limitation has led to the development of various workarounds, such as sub-interpreters, multiprocessing, and C extensions, to circumvent the GIL and improve performance - or even remove it entirely, as proposed in PEP 703⁴, which was accepted in Python 3.13 and is currently in the experimental stage. This project aims to investigate the performance of GIL-free Python for compute-bound tasks and compare it to alternative approaches.

What’s a GIL?

Motivation:

- Memory/Network-bound tasks: Asynchronous I/O with `asyncio`, very competitive.
- Compute-bound tasks: Very slow interpreter, hard to parallelize with GIL. → recently removed in PEP 703

Research question:

- How useful is GIL-free Python for compute-bound tasks?
- How does it compare to alternatives (multiprocessing, C-Python interopt, C-Python extensions)?

Chosen algorithm: hashcat

- on password storage: https://cheatsheetseries.owasp.org/cheatsheets/Password_Storage_Cheat_Sheet.html
- we use a simpler one
- no algorithmic optimizations (e.g. rainbow tables, bloom filters, etc.) just brute-force

Cpython dependency `Python.h`: <https://github.com/python/cpython/blob/main/Include/Python.h>

Experiments




Figure 1: Performance Overview

command	mean	stddev	median	user	system	min	max
plain: itertools.py	0.5674692	0.0119883	0.5700655	0.5599028	0.0074184	0.5496380	0.5865690
plain: lib.py	0.1003698	0.0026018	0.0995592	0.0950988	0.0051505	0.0976092	0.1086287
plain: plain.py	0.5631182	0.0085463	0.5607433	0.5574100	0.0056163	0.5564152	0.5863659
multiprocessing: imap_unordered.py	0.2258019	0.0085966	0.2235880	0.5456730	0.1661910	0.2184310	0.2449692
multiprocessing: imap.py	0.2328316	0.0065632	0.2306093	0.5554529	0.1625106	0.2235672	0.2426373
multiprocessing: map_async.py	0.4528332	0.0248580	0.4485649	1.0100400	0.1108107	0.4314743	0.5167882
multiprocessing: map.py	0.4400746	0.0043315	0.4405771	0.9853628	0.1084339	0.4329375	0.4467715
multithreading: GIL=1 executor.py	0.3696592	0.0103798	0.3658508	0.3597924	0.0092238	0.3621005	0.3968615

¹<https://github.blog/news-insights/octoverse/octoverse-2024/#the-most-popular-programming-languages>
²Wang, Z., Bu, D., Sun, A., Gou, S., Wang, Y., & Chen, L. (2022). An empirical study on bugs in python interpreters. IEEE Transactions on Reliability, 71(2), 716-734.
³<https://benchmarksgame-team.pages.debian.net/benchmarksgame/index.html>
⁴<https://peps.python.org/pep-0703/>

command	mean	stddev	median	user	system	min	max
multithreading: GIL=0 executor.py	0.2102704	0.0104267	0.2121045	0.4389796	0.0105094	0.1962787	0.2321854
multithreading: GIL=1 workers.py	0.1304648	0.0071726	0.1292655	0.1178397	0.0081775	0.1242261	0.1582329
multithreading: GIL=0 workers.py	0.1677349	0.0076839	0.1685633	0.1931002	0.0202853	0.1429964	0.1760283
ctypes: invoke_hashcat.py	0.0934947	0.0031416	0.0929726	0.0882496	0.0049164	0.0891827	0.0996272
ctypes: invoke_hashcat.py (openmp)	0.1021338	0.0056378	0.1003631	0.0986943	0.0083828	0.0976012	0.1269725
cpython: invoke_hashcat.py	0.1006056	0.0043579	0.0997623	0.0950439	0.0052310	0.0943297	0.1081794

Target hash: aaa

Warmup: 3 runs

Docker with Python 3.13t experimental build

Addendum

The following system specifications were used for the experiments:

```
$ system_profiler SPSoftwareDataType SPHardwareDataType
```

Software:

```
System Software Overview:

    System Version: macOS 14.6.1 (23G93)
    Kernel Version: Darwin 23.6.0
    Boot Volume: Macintosh HD
    Boot Mode: Normal
    Computer Name: Yahya’s MacBook Pro
    User Name: Yahya Jabary (sueszli)
    Secure Virtual Memory: Enabled
    System Integrity Protection: Enabled
    Time since boot: 79 days, 22 hours, 26 minutes
```

Hardware:

```
Hardware Overview:

    Model Name: MacBook Pro
    Model Identifier: Mac14,10
    Model Number: Z174001ABD/A
    Chip: Apple M2 Pro
    Total Number of Cores: 12 (8 performance and 4 efficiency)
    Memory: 16 GB
    System Firmware Version: 10151.140.19
    OS Loader Version: 10151.140.19
    Serial Number (system): VCYQDOHHOG
    Hardware UUID: BEA4D09D-6651-54E1-A3F7-7FB78A7BF1AB
    Provisioning UDID: 00006020-001A284901E8C01E
    Activation Lock Status: Disabled
```

```
$ docker compose exec main lscpu
Architecture:           x86_64
CPU op-mode(s):         32-bit
Byte Order:             Little Endian
CPU(s):                 12
On-line CPU(s) list:    0-11
Thread(s) per core:     1
Core(s) per socket:     12
Socket(s):              1
Vendor ID:              0x61
Model:                  0
Stepping:               0x0
BogoMIPS:               48.00
Vulnerability Gather data sampling: Not affected
Vulnerability Itlb multihit: Not affected
Vulnerability L1tf:     Not affected
Vulnerability Mds:      Not affected
Vulnerability Meltdown: Not affected
Vulnerability Mmio stale data: Not affected
Vulnerability Reg file data sampling: Not affected
Vulnerability Retbleed: Not affected
Vulnerability Spec rstack overflow: Not affected
Vulnerability Spec store bypass: Mitigation; Speculative Store Bypass disabled via prctl
Vulnerability Spectre v1: Mitigation; __user pointer sanitization
Vulnerability Spectre v2: Not affected
Vulnerability Srbds:     Not affected
Vulnerability Tsx async abort: Not affected
Flags:                  fp asimd evtstrm aes pmull sha1 sha2 crc32 atomics fphp asimdhp cpuid asimdrdm jscvt fcma lrcpc dcpop
                        ilrcpc flagm ssbs sb paca pacg dcpodp flagm2 frint
```